Does the Grey Lady treat everyone equally?

Quinn Underriner

Introduction:

This era is marked by heightened scrutiny of the authority of news sources. As distribution methods of information have democratized through the proliferation of social media, there is a heightened cultural debate about what constitutes objectivity and truth. This inquiry is not just for the new media newcomers but also for venerable institutions such as the New York Times ("Fake News", as called by one popular public personality). In response, the New York Times, the paper of record, the liberal news source with the largest readership and highest number of Pulitzer Prizes, has positioned itself as a beacon of truth and the voice for Liberalism in this time of uncertainty.

This paper examines over two thousand articles articles scraped from the New York

Times website between April and October 2019 related to the 2020 Democratic primary and
conducts two primary forms of analysis:

- i. Descriptive natural language processing highlighting how and with what frequency the New York Times has covered democratic nominees in the primary
- ii. Study if there are any major discrepancy between (as measured through polling and ANES data) the popularity of candidates in the general population of likely Democratic Party voters and how favorably they are discussed in the New York Times, and if a discrepancy exists, to see if this bias happens on ideological lines. While the New York Times has no obligation to mirror the population, it is useful to see how it deviates from the average "liberal", here loosely defined as the average Democratic Party voter.

¹ "Pulitzer Prizes," The New York Times Company. Retrieved December 3, 2019.

The bias element of the second claim is one that has been leveled at the media writ large by the Bernie Sanders campaign. Nina Turner, the Co-National Chair for the Presidential Campaign of Bernie Sanders, recently was interviewed saying "the Bernie blackout is real, its not a figment of our imagination... even though he is polling very high, either number or number 2 [in the polls]... these networks are wired to only care about the upper middle class and the ultra wealthy". She goes on to note that he has raised more money from more individuals than any other candidate and reached one million individual donors before any other candidate this September. To examine this claim this paper will create a measure of ideology among the Democratic nominee hopefuls.

Sources of Data:

I used the Selenium and Beautiful Soup packages to scrape the New York Times website for the phrase "2020 election" across all of their articles (from all categories: politics, op-ed, business, etc.) between May and October of 2019. I gathered 2,100 articles.

search_term	title	date	url	text
2020 Election	Do More Candidates Frighten You?	Oct. 31	https://www.nytimes.com/2019/10/31/us/politics	On Politics We asked, you answered: Is there a

This resulted in a data-frame that looked like the above, with the article headline, the date of publication, the url, and the full text as features. There is some risk that coverage of the election will be missed if an article does not contain the phrase "2020 Election", but with 2,100

² "Nina Turner talks Sanders media coverage", The Hill. Retrieved December 10, 2019, https://thehill.com/hilltv/rising/473191-nina-turner-talks-sanders-media-coverage

³ "2020 Presidential Race", Center for Responsive Politics, Retrieved December 9, 2019 https://www.opensecrets.org/2020-presidential-race

articles there should be enough material to capture editorial preferences. Further, while this

process might pick up articles that are only tangentially related to the election, as discussed

below, the final text used in analysis is further filtered for mentions of specific candidates, so

ultimately its likely irrelevant article will not be used. While of course the New York Times will

employee writers across the ideological spectrum, I believe the size of this sample will be able to

capture if writers of a certain bent are more frequently given assignments.

The articles' text was preprocessed to make all the text lowercase, and stripped of

punctuation and stopwords (these are words like "the" that do not add to our understating of the

underlying corpus). I then used what is the most common stemming procedure, Porter's

algorithm, to produce a list of stemmed words. This algorithm works to remove suffixes to

standardize related words and return a comparable root⁴ (e.g., "house", "houses" and "housing"

would all become "hous"; they are all referring to the same object and should be treated as such).

After all of this preprocessing has occurred there are roughly one and a half million words to be

analyzed.

To map words to candidates I took the candidates last name (being careful to search for

this name in its stemmed form) and then stored the ten words before and after this keyword in a

nested list.

Across the corpus there were the following number of mentions per candidate:

Biden: 4797 mentions

Warren: 2287 mentions

Bernie: 1465 mentions

⁴ Christopher D. Manning, Prabhakar Raghavan and Hinrich Schütze, Introduction to Information Retrieval, Cambridge University Press. 2008. Retrieved December 5th, 2019, https://nlp.stanford.edu/IRbook/html/htmledition/stemming-and-lemmatization-1.html

Harris: 1026

Buttigieg: 660 mentions

Klobuchar: 533

Beto: 368 mentions

These mentions were further broken down by month so that they could be mapped to polling data.

For polling data, I used data from Real Clear Politics (RCP),⁵ a polling aggregator, to find monthly average polling for each candidate across a variety of polls. RCP relies more heavily on some sources that others, with the largest number of polls coming from Politico/Morning Consult and Economist/YouGov, as seen in the chart below.

Politico/Morning Consult	33
Economist/YouGov	26
The Hill/HarrisX	16
Quinnipiac	12
Emerson	8
Harvard-Harris	7
CNN	7
FOX News	7
Monmouth	5
SurveyUSA	4
IBD/TIPP	4
USA Today/Suffolk	3
NBC News/Wall St. Jrnl	3
Reuters/Ipsos	3
LA Times/USC	3
ABC News/Wash Post	3
RCP Average	1

The polls are all taken over a period of roughly 3-5 days. I took the starting month to assign each poll to a given month. This means that there are some polls that straddled months that were assigned to month that they started in.

RCP stops collecting polling for candidates when they drop out of the race, naturally. This meant that while I was able to grab data for Kamala Harris before she left the running, Beto

⁵ "Latest 2020 Democratic Presidential Primary Polls", Real Clear Politics, Accessed December 5th, 2019 from https://www.realclearpolitics.com/epolls/latest polls/democratic nomination polls/

O'Rourke had already dropped out of the data set before I scrapped it using a command line tool.⁶

Below is the monthly ranking from each poll sorted by the October rank. Biden is clearly leading across all six months, with Sanders and Warren swapping second and third in September. Klobuchar is notably in perpetual last place.

Candidate	May_Poll_Rank	June_Poll_Rank	July_Poll_Rank	August_Poll_Rank	Sep_Poll_Rank	Oct_Poll_Rank
Biden	1.0	1.0	1.0	1.0	1.0	1.0
Warren	3.0	3.0	3.0	3.0	2.0	2.0
Sanders	2.0	2.0	2.0	2.0	3.0	3.0
Buttigieg	5.0	5.0	5.0	5.0	5.0	4.0
Harris	4.0	4.0	4.0	4.0	4.0	5.0
Klobuchar	6.0	6.0	6.0	6.0	6.0	6.0

The final source of data was the 2018 Pilot Study from the American National Election Studies (ANES).⁷ I used the pilot study as the comprehensive study for 2018 has not yet been released (this was the latest data available). The main difference between the pilot study and the full study is the sample size, N=2,500, is about half as large. This survey gathers detailed demographic and political opinion data from a statistically representative portion of Americans. All participants are US citizens over the age of 18.

While there is a wealth of information in the ANES, I used two features in this analysis. First, "ideo5", which is a 5 point (1 being the most liberal, 5 being the most conservative, with a 6th option for "unsure") scale of voters' self-identified political preferences. Second, "vote20cand" which asks "In the 2020 Democratic primary for president, who will you vote for?

⁶ Anthony Bloomer, "Real Clear Politics," https://pypi.org/project/realclearpolitics/

⁷ "ANES 2018 Pilot Study Questionnaire Specifications", American National Election Studies, Retrieved December 1, 2019, https://electionstudies.org/wp-content/uploads/2018/12/ anes pilot 2018 questionnaire.pdf

Your best guess is fine" and gives 11 options. Of the 2500 people in the survey, 1243 of the 2500 answered this question (presumably some people didn't know and a large number are Republicans). Below is the sum of results from "vote20cand". Compared the poll results, which reflect a later point in time, we see similarity at both the top (with Biden at a comfortable lead and Bernie in second) and at the bottom (Klobuchar), but otherwise there is some shuffling

(notably the ascendency of Warren in the polls).

votezocano
391
246
183
115
105
54
49

It's important to note that some later entrants to the presidential race, like Pete Buttigieg, are absent from this list (he announced his candidacy April 14, 2019) and therefore is not included in the analysis that involves ANES data. This data is used together to gather a representative sample of popularity of each given candidate in the general population, as well as place candidates on an ideological spectrum. To do this latter analysis it is assumed that who voters say they are voting for is the candidate which most aligns with their values. This seems like a fair assumption, but voters could be factoring in things like "electability" over their own policy preferences. They could also be poorly informed, either about a candidates actual policies or where their own politics land on the US political spectrum. This could slightly skew the results. A final caveat is that this survey was administered between December 5 and December 21 of 2018. New information could have come to light about candidates after this time period

that would have changed peoples opinions by the time roughly six months later that the New York Times sentiment analysis started.

Analysis:

Below are two word clouds, made from the entire six month corpus.



Bernie Sanders



Joe Biden

Word clouds can be useful for some initial exploratory analysis of our data. We can see that the scandal involving Hunter Biden on the board of a Ukrainian gas company has dominated Joe Biden's coverage in the New York Times (which is reiterated later in this paper when topic modeling is applied to his corpus), while coverage of Bernie Sanders has a more of a policy

focus, especially relating to healthcare, as well as a focus on his number of donors. Simple analysis at this level does not bear out significant bias.

Below are the aggregate sentiment scores for each candidate in their New York Times coverage. These are calculated by creating a list of all the words associated with each person's name and adding up the sentiment score for each word. I used the Affin method to calculate sentiment which takes each word fed to it and returns a numerical score from -5 to 5 based on its perceived positivity or negativity. This is an admittedly naive approach as word can have different meanings in the aggregate than they do alone, but I'm hoping the large size of the sample will smooth out any minor misclassifications.

	sentiment_scores	sentiment_rank
Warren	2262.0	1.0
Biden	1491.0	2.0
Sanders	1382.0	3.0
klobuchar	678.0	4.0
beto	448.0	5.0
harris	417.0	6.0
Buttigleg	330.0	7.0

I used Afinn over Bing, another popular method, because I wanted more granularity in the scores so I would be more likely to register differences between the candidates. Bing just produces a score of -1 (for negative), 0 (for neutral), or 1 (for positive). We can see that by magnitude Warren has had outsize positive coverage. Klobuchar, who is trailing in the polls and the ANES survey, is doing relatively very well. Coverage of Sanders is decently positive.

In the below chart I'm simply calculating the difference between this sentiment rank above and the rank given by the stated voting preference of participants in the ANES survey (rank being created by the number of people intending to vote for that candidate) and I'm then subtracting the sentiment rank from the voting rank. The difference column can be read as the higher the positive value of the difference column number, the relatively lower the coverage received from the NYTimes from the public opinion baseline. Klobuchar and Warren are clear winners, Harris is the largest loser, and Biden and Sanders both have equally slightly less positive than expected coverage.

	Candidate	Sentiment Rank	Voting Rank	Difference
0	Joe Biden	2	1	1.0
1	Bernie Sanders	3	2	1.0
2	Kamala Harris	6	4	2.0
3	Elizabeth Warren	1	5	-4.0
4	Amy Klobuchar	4	7	-3.0

Below is the same analysis, but done at a monthly level using the aggregated Real Clear Politics Polling data. The way to read the "Coverage Gap" column is that if the number is negative the New York Times covers them positively relatively less than their popularity among voters and if the number is positive then the New York Times covers them relatively more glowingly than their popularity among voters.

	diff_oct	diff_sep	diff_may	diff_june	diff_july	diff_aug	Coverage Gap
Biden	-2.0	-2.0	0.0	-1.0	-3.0	-2.0	-10.0
Harris	-1.0	-3.0	0.0	-2.0	0.0	-1.0	-7.0
Buttigieg	-1.0	-0.5	-1.0	0.0	1.0	-2.0	-3.5
Sanders	1.0	1.0	0.0	-2.0	-2.0	0.0	-2.0
Warren	1.0	1.0	0.0	2.0	-1.0	2.0	5.0
Klobuchar	2.0	2.0	1.0	3.0	2.0	2.0	12.0

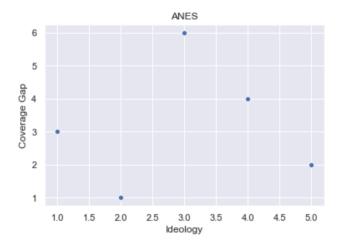
Biden got more overall mentions from the New York Times than the next two largest number of mentions (for Warren and Sanders, respectively) combined. While he's getting a lot of press it's clearly not entirely positive, as shown by his high score; topic model analysis can provide some insight into this. If you set the parameters of the topic model to find just three in the corpus one of the three is related to the corruption allegations involving his son Hunter. Likely much of the negative sentiment comes from related press. It is important to note that as Trump is involved in this scandal as well, this process could also be picking up New York Times' ire directed at Trump and artificially lowering Biden's score.

To gauge the ideology of the candidates, I ran regressions on dummy variables for each of the six self-described points on the ideological spectrum in the ANES. For example, for Bernie Sanders, both the most liberal feature (which had the highest coefficient of any of the candidates) as well as the "I don't know" dummies are the only ones that are statistically significant (see chart below), but for Joe Biden while the most liberal dummy variable is statistically significant, so is the 3rd and 4th point (indicating moderate to conservative leanings) on the spectrum.

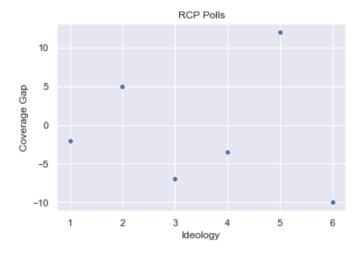
	coef	std err	t	P> t	[0.025	0.975]
const	0.1782	0.014	13.122	0.000	0.152	0.205
ideo5_1	0.0579	0.024	2.414	0.016	0.011	0.105
ideo5_2	-0.0166	0.022	-0.745	0.456	-0.060	0.027
ideo5_3	0.0088	0.022	0.404	0.686	-0.034	0.052
ideo5_4	-0.0115	0.038	-0.304	0.761	-0.086	0.063
ideo5_5	-0.0032	0.055	-0.058	0.954	-0.111	0.104
ideo5 6	0.1428	0.040	3.598	0.000	0.065	0.221

I rank ordered the candidates based on the statistical significance of these relative dummy variables, as well as the magnitudes of the related coefficients, with higher coefficients and more significant lower numbered dummies making a candidate more "liberal". At the end of this analysis, my ranking of candidates from least liberal to most liberal is: Biden (1), Klobuchar (2), Beto (3), Harris (4), Warren (5) and then Bernie (6).

Below is a scatter plot of these ideology scores vs. the coverage gap figure from above using the ANES rankings.



Doing the same analysis with the RCP polling data:



In both graphs the N is extremely small, but there does seem to be a slight up and to the right trend in both cases (with Biden being a strong outlier, again likely due to the Ukraine scandal), indicating, generally and weakly, that the coverage gap increases as a candidate becomes more liberal.

Ultimately, in response to Nina Turner charge, Senator Sanders does receive significantly less overall coverage than his placement in either the polls or in the ANES rankings. However, the sentiment of the coverage he does receive is close to commensurate with his popularity. The main beneficiary of any alleged New York Times preference for a more centrist candidate is Amy Klobuchar who receives both more total coverage, and especially outsized positive coverage, given the low measures of her popularity. Elizabeth Warren, who is notably the second most leftward candidate according to the above analysis, also has a more pronounced increase in overall coverage and a moderate boost in sentiment. Further and more robust analysis is certainly called for, but these initial results are far from damning or indicative of a strong bias.

```
In [112]: import pandas as pd
          import numpy as np
          import nltk
          import unicodedata
          import sys
          import re
          import os
          from nltk.corpus import stopwords
          from nltk.tokenize import word_tokenize
          from wordcloud import WordCloud, STOPWORDS, ImageColorGenerator
          from sklearn.feature_extraction.text import CountVectorizer
          from sklearn.decomposition import LatentDirichletAllocation as LDA
          from nltk.stem import PorterStemmer
          from afinn import Afinn
          import seaborn as sns
          from PIL import Image
          import matplotlib.pyplot as plt
```

```
In [288]: | # The below data scraped the websites
          import time
          import requests
          import pandas as pd
          from bs4 import BeautifulSoup
          import json
          import string
          from selenium import webdriver
          from selenium.webdriver.common.by import By
          from selenium.webdriver.support.ui import WebDriverWait
          from selenium.webdriver.support import expected conditions as EC
          #from webdriver manager.chrome import ChromeDriverManager
          import os
          def get_articles(search_term, start_date, end_date):
              pwd = os.getcwd()
              base = "https://www.nytimes.com"
              browser = webdriver.Chrome('/Users/quinnunderriner/Desktop/school/Co
          mp American/nytimes v2 REAL/chromedriver')
              wait = WebDriverWait(browser, 10)
              search_term_url = search_term.replace(' ', '%20')
              search_url = 'https://www.nytimes.com/search?endDate={}&query={}&sor
          t=newest&startDate={}'.format(end date,
          search term url,
          start date)
              browser.get(search url)
              while True:
                  try:
                       time.sleep(1)
                       show_more = wait.until(EC.element_to_be_clickable((By.XPATH,
          '//button[@type="button"][contains(., "Show More")]')))
                       show more.click()
                  except Exception as e:
                           print(e)
                          break
              soup = BeautifulSoup(browser.page source, 'lxml')
              search results = soup.find('ol', {'data-testid':'search-results'})
              articles = []
              links = search results.find all('a')
              for link in links:
                  link url = link['href']
                  try:
                      title = link.find('h4').text
                      date = link.find next('time').text
                      print(date + ': '+ title + '\n')
                       response = requests.get(base + link url)
                       soup link = BeautifulSoup(response.text, 'html.parser')
```

```
scripts = soup link.find all('script')
            for script in scripts:
                if 'window. preloadedData = ' in script.text:
                    jsonStr = script.text
                    jsonStr = jsonStr.split('window. preloadedData = ')
[-1]
                    jsonStr = jsonStr.rsplit(';',1)[0]
                    jsonData = json.loads(jsonStr)
                    article = []
                    for k, v in jsonData['initialState'].items():
                        w=1
                        try:
                            if v[' typename'] == 'TextInline':
                                article.append(v['text'])
                                #print (v['text'])
                        except:
                            continue
                    article = [ each.strip() for each in article ]
                    article = ''.join([('' if c in string.punctuation el
se ' ')+c for c in article]).strip()
            articles.append([search term, title, date, base+link url, ar
ticle])
        except:
            continue
    print("Complete")
    df = pd.DataFrame(articles, columns=['search term', 'title', 'date',
'url', 'text'])
    df.to_csv('{}_{{}}_{{}}.csv'.format(search_term.replace(' ', '_'), start
date, end date), index=False)
    browser.quit()
    return df
#Called the scraping function this way, it writes the output of each bat
ch to csvs and I put each piece together below
df = get articles('2020 Election', '20191016', '20191031')
```

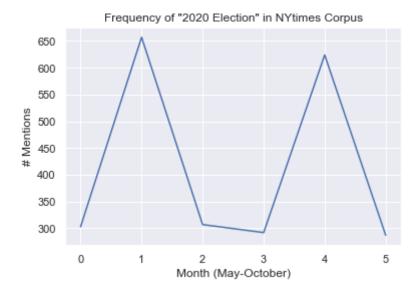
```
In [ ]: # the below function process text
        def keep_chr(char):
            return (unicodedata.category(char).startswith('P'))
        PUNCTUATION = " ".join(
            [chr(i) for i in range(sys.maxunicode) if keep_chr(chr(i))])
        stop_words = set(stopwords.words('english'))
        def preprocess(df):
            df = df.lower()
            df = re.sub(r'\d+', '', df)
            porter = PorterStemmer()
            clean_df = []
            df_split = df.split()
            for i in df split:
                i = i.strip(PUNCTUATION)
                if i in stop words:
                     continue
                if len(i) == 0:
                    continue
                i = porter.stem(i)
                clean_df.append(i)
            return clean_df
        def clean row level(x):
            return preprocess(x["text"])
```

```
In [3]: df["cleaned"] = df.apply(lambda x: clean_row_level(x), axis=1)
#after cleaning and writing to csv using the above functions, this loads
    the saved full corpus
    df = pd.read_csv("latest_full_articles.csv")
```

```
In [ ]: #this is the code that creates wordclouds
        def create cloud candidates(df):
             # https://www.datacamp.com/community/tutorials/wordcloud-python
            stop_words = set(stopwords.words('english'))
            #add specific stopwords plus names of candidates
            stop words.update(["would","said","mr","presid","democrat","new","li
        ke","elect","year","mrs", \
                                "american", "also", "could", "whether", "day", "go", "us
        e", "two", "one", "campaign", \
                                "call", "even", "say", "get", "may", "make", "come", "cam
        paign","state","polit", \
                                "time", "ms", "call", "elizabeth", "warren", "joe", "bid
        en","pete","buttigieg","bernie", \
                                "sander", "berni", "debat", "candid", "first", "second"
        ,"percent","night","poll", \
                                "vice", "vermont", "massachusett", "joseph", "kamala",
        "south", "bend", "senat", "harri", "race", "former", "th"])
            text = " ".join(df)
            usa_mask = np.array(Image.open("US_img.png"))
            wordcloud = WordCloud(max_words=100, stopwords=stop_words, \
                         background_color="white", collocations=False,\
                         mask=usa mask, contour width=2, contour color='steelblu
        e').generate(text)
            plt.figure(figsize=(25,20))
            plt.imshow(wordcloud, interpolation="bilinear")
            plt.axis("off")
            plt.show()
```

```
In [362]: import numpy as np
    y = [303,657,307,292,624,287]
    x = np.arange(0,len(y))
    import seaborn as sns; sns.set()
    import matplotlib.pyplot as plt
    ax = sns.lineplot(x=x, y=y)
    plt.title('Frequency of "2020 Election" in NYtimes Corpus')
    plt.xlabel('Month (May-October)')
    plt.ylabel('# Mentions')
```

Out[362]: Text(0, 0.5, '# Mentions')



```
In [4]: #standarize the dates
def strip_date(x):
    return x["date"][0:3]
df["date"] = df.apply(lambda x: strip_date(x), axis=1)
```

```
In [5]: #break out data by month to be able to do month to month comparisons.
    df_may = df[df["date"] == "May"]
    df_june = df[df["date"] == "Jun"]
    df_jul = df[df["date"] == "July"]
    df_aug = df[df["date"] == "Aug"]
    df_sep = df[df["date"] == "Sep"]
    df_oct = df[df["date"] == "Oct"]
```

```
In [24]: def gimmie month rank(df, month):
             returns a dataframe with the sentimetn scores for each candidate in
          a given month
             df = df.text.str.cat()
             df = preprocess(df)
             sanders list, warren list, biden list, buttigieg list = extact relat
         ed words(df)
             sanders_list_text = " ".join(map(str, sanders_list))
             warren_list_text = " ".join(map(str, warren_list))
             biden list_text = " ".join(map(str, biden_list))
             buttigleg list text = " ".join(map(str, buttigleg list))
             afinn = Afinn()
             biden score = afinn.score(biden list text)
             sanders score = afinn.score(sanders list text)
             buttigieg_score = afinn.score(buttigieg_list_text)
             warren_score = afinn.score(warren list text)
             df2 = pd.DataFrame(np.array([[biden score],
                                            [sanders score],
                                            [buttigieg score],
                                            [warren_score]]), columns=[month + "_s
         entiment_score"],index=["Biden","Sanders","Warren","Buttigieg"])
             return df2
```

```
In [33]: #this loop creates a dataframe with the sentiment scores for each month
    month_list = ["june","jul","aug","sep","oct"]
    start = ny.gimmie_month_rank(df_may,"may")
    start = start.rank(axis=0, ascending=False)
    for i, j in enumerate([df_june, df_jul, df_aug, df_sep, df_oct]):
        new_month = ny.gimmie_month_rank(j,month_list[i])
        new_month = new_month.rank(axis=0, ascending=False)
        start = pd.concat([start, new_month], axis = 1)
```

```
In [34]: start
```

Out[34]:

	may_sentiment_score	june_sentiment_score	jul_sentiment_score	aug_sentiment_score
Biden	1.0	2.0	4.0	3.0
Sanders	2.0	4.0	4.0	2.0
Buttigieg	6.0	5.0	4.0	7.0
Warren	3.0	1.0	4.0	1.0
Beto	7.0	7.0	4.0	6.0
Harris	4.0	6.0	4.0	5.0
Klobuchar	5.0	3.0	4.0	4.0

```
In [18]: def generate ranked polls(df):
             This merges RCP polls and creates ranked lists from them
             candidate list = ["Biden", "Sanders", "Warren", "Buttigieg", "Harris", "K
         lobuchar","Date"]
             def strip date2(x):
                 return x["Date"][0:2]
             #polling data grabbed from real clear politics using https://pypi.or
         g/project/realclearpolitics/
             #command line tool
             #polls = pd.read csv("rcp polls PLUS.csv")
             df["Date"] = df.apply(lambda x: strip date2(x), axis=1)
             df = df[candidate list]
             print(df.columns)
             df = df.groupby("Date").mean()
             df = df.T
             df = df \cdot drop(["4/"], axis=1)
             df["Oct_Poll_Rank"] = df["10"].rank(ascending=False)
             df["Nov Poll Rank"] = df["11"].rank(ascending=False)
             df["May_Poll_Rank"] = df["5/"].rank(ascending=False)
             df["June Poll Rank"] = df["6/"].rank(ascending=False)
             df["July Poll Rank"] = df["7/"].rank(ascending=False)
             df["August_Poll_Rank"] = df["8/"].rank(ascending=False)
             df["Sep Poll Rank"] = df["9/"].rank(ascending=False)
             df = df[['May Poll Rank', 'June Poll Rank', 'July Poll Rank',
                 'August_Poll_Rank', 'Sep_Poll_Rank','Oct_Poll_Rank']]
             return df
```

```
In [9]: #read in polling data
polls = pd.read_csv("rcp_polls_PLUS.csv")
```

```
In [19]: #create poll rank df
         six candidate rank polls = generate ranked polls(polls)
```

Index(['Biden', 'Sanders', 'Warren', 'Buttigieg', 'Harris', 'Klobucha r', 'Date'], dtype='object')

In [397]: #Concat polling data and sentiment data

final nytime sentiment diff = pd.concat([six candidate rank polls, start],axis=1)final nytime sentiment diff = final nytime sentiment diff.dropna(axis=0)

/Users/quinnunderriner/anaconda3/lib/python3.7/site-packages/ipykernel launcher.py:2: FutureWarning: Sorting because non-concatenation axis is not aligned. A future version of pandas will change to not sort by default.

To accept the future behavior, pass 'sort=False'.

To retain the current behavior and silence the warning, pass 'sort=Tru e'.

In [36]: final nytime sentiment diff

Out[36]:

	May_Poll_Rank	June_Poll_Rank	July_Poll_Rank	August_Poll_Rank	Sep_Poll_Rank	Oc
Biden	1.0	1.0	1.0	1.0	1.0	
Buttigieg	5.0	5.0	5.0	5.0	5.0	
Harris	4.0	4.0	4.0	4.0	4.0	
Klobuchar	6.0	6.0	6.0	6.0	6.0	
Sanders	2.0	2.0	2.0	2.0	3.0	
Warren	3.0	3.0	3.0	3.0	2.0	

In [290]: #calculate the difference between nytimes sentiment and polling data for each month final_nytime_sentiment_diff["diff_oct"] = final_nytime_sentiment_diff["o

ct sentiment score"] - final nytime sentiment diff["Oct Poll Rank"] final nytime sentiment diff["diff sep"] = final nytime sentiment diff["s ep sentiment score"] - final nytime sentiment diff["Sep Poll Rank"] final nytime sentiment diff["diff may"] = final nytime sentiment diff["m ay sentiment score"] - final nytime sentiment diff["May Poll Rank"] final_nytime_sentiment_diff["diff_june"] = final nytime sentiment diff["june sentiment score" | - final nytime sentiment diff["June Poll Rank"] final nytime sentiment diff["diff july"] = final nytime sentiment diff["jul sentiment score"] - final nytime sentiment diff["July Poll Rank"] final_nytime_sentiment_diff["diff_aug"] = final_nytime_sentiment_diff["a ug sentiment score"] - final nytime sentiment diff["August Poll Rank"]

```
In [396]: | final_nytime_sentiment_diff
```

Out[396]:

	diff_oct	diff_sep	diff_may	diff_june	diff_july	diff_aug	sum
Biden	2.0	2.0	0.0	1.0	3.0	2.0	10.0
Buttigieg	1.0	0.5	1.0	0.0	-1.0	2.0	3.5
Harris	1.0	3.0	0.0	2.0	0.0	1.0	7.0
Klobuchar	-2.0	-2.0	-1.0	-3.0	-2.0	-2.0	-12.0
Sanders	-1.0	-1.0	0.0	2.0	2.0	0.0	2.0
Warren	-1.0	-1.0	0.0	-2.0	1.0	-2.0	-5.0

In [398]: #CORRECT ONE

#calculate the difference between nytimes sentiment and polling data for each month

final_nytime_sentiment_diff["diff_oct"] = final_nytime_sentiment_diff[
"Oct_Poll_Rank"] - final_nytime_sentiment_diff["oct_sentiment_score"]
final_nytime_sentiment_diff["diff_sep"] = final_nytime_sentiment_diff[
"Sep_Poll_Rank"] - final_nytime_sentiment_diff["sep_sentiment_score"]
final_nytime_sentiment_diff["diff_may"] = final_nytime_sentiment_diff["M
ay_Poll_Rank"] - final_nytime_sentiment_diff["may_sentiment_score"]
final_nytime_sentiment_diff["diff_june"] = final_nytime_sentiment_diff[
"June_Poll_Rank"] - final_nytime_sentiment_diff["june_sentiment_score"]
final_nytime_sentiment_diff["diff_july"] = final_nytime_sentiment_diff[
"July_Poll_Rank"] - final_nytime_sentiment_diff["jul_sentiment_score"]
final_nytime_sentiment_diff["diff_aug"] = final_nytime_sentiment_diff["A
ugust_Poll_Rank"] - final_nytime_sentiment_diff["aug_sentiment_score"]

In [399]: sentiment_diff = final_nytime_sentiment_diff[["diff_oct","diff_may","dif
f_sep","diff_july","diff_june","diff_aug"]]
sentiment_diff = sentiment_diff.abs()
sentiment_diff.sum().sum()/36
#Nytimes is on average deviates by 1.3 from the polling data

Out[399]: 1.3194444444444444

/Users/quinnunderriner/anaconda3/lib/python3.7/site-packages/ipykernel_launcher.py:3: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame. Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy

This is separate from the ipykernel package so we can avoid doing imports until

In [402]: final_nytime_sentiment_diff.sort_values(by=['sum'])

Out[402]:

_		diff_oct	diff_sep	diff_may	diff_june	diff_july	diff_aug	sum
-	Biden	-2.0	-2.0	0.0	-1.0	-3.0	-2.0	-10.0
	Harris	-1.0	-3.0	0.0	-2.0	0.0	-1.0	-7.0
	Buttigieg	-1.0	-0.5	-1.0	0.0	1.0	-2.0	-3.5
	Sanders	1.0	1.0	0.0	-2.0	-2.0	0.0	-2.0
	Warren	1.0	1.0	0.0	2.0	-1.0	2.0	5.0
	Klobuchar	2.0	2.0	1.0	3.0	2.0	2.0	12.0

```
In [427]: sorted_sentiment = final_nytime_sentiment_diff.sort_values(by=['sum'])
```

```
In [446]: ideology = np.array([6,3,4,1,2,5])
sorted_sentiment['ideology'] = np.array(ideology)
```

In [447]: sorted_sentiment

Out[447]:

	diff_oct	diff_sep	diff_may	diff_june	diff_july	diff_aug	sum	ideology
Biden	-2.0	-2.0	0.0	-1.0	-3.0	-2.0	-10.0	6
Harris	-1.0	-3.0	0.0	-2.0	0.0	-1.0	-7.0	3
Buttigieg	-1.0	-0.5	-1.0	0.0	1.0	-2.0	-3.5	4
Sanders	1.0	1.0	0.0	-2.0	-2.0	0.0	-2.0	1
Warren	1.0	1.0	0.0	2.0	-1.0	2.0	5.0	2
Klobuchar	2.0	2.0	1.0	3.0	2.0	2.0	12.0	5

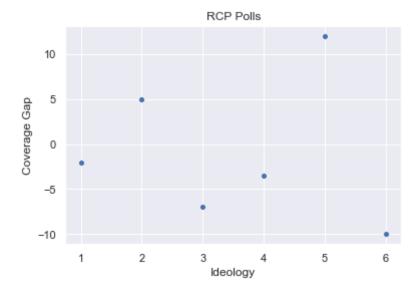
In [456]: sorted sentiment.rename(columns={"sum": "Coverage Gap"})

Out[456]:

	diff_oct	diff_sep	diff_may	diff_june	diff_july	diff_aug	Coverage Gap	ideology
Biden	-2.0	-2.0	0.0	-1.0	-3.0	-2.0	-10.0	6
Harris	-1.0	-3.0	0.0	-2.0	0.0	-1.0	-7.0	3
Buttigieg	-1.0	-0.5	-1.0	0.0	1.0	-2.0	-3.5	4
Sanders	1.0	1.0	0.0	-2.0	-2.0	0.0	-2.0	1
Warren	1.0	1.0	0.0	2.0	-1.0	2.0	5.0	2
Klobuchar	2.0	2.0	1.0	3.0	2.0	2.0	12.0	5

```
In [459]: ax = sns.scatterplot(x="ideology", y="sum", data=sorted_sentiment)
    plt.ylabel('Coverage Gap')
    plt.xlabel('Ideology')
    plt.title('RCP Polls')
```

Out[459]: Text(0.5, 1.0, 'RCP Polls')



```
In [ ]:
```

```
In [88]: #load anes data
anes = pd.read_csv("anes_pilot_2018.csv")
```

/Users/quinnunderriner/anaconda3/lib/python3.7/site-packages/IPython/core/interactiveshell.py:2785: DtypeWarning: Columns (69,78) have mixed types. Specify dtype option on import or set low_memory=False. interactivity=interactivity, compiler=compiler, result=result)

```
In [90]: #drop unneeded columns from ANES and apply cadidate name dictionary gene
         rated above
         anes_sub_keep = anes[["vote20cand","pid1r","media1","media2","media3","m
         edia4", "trustmedia", "gender", "race", "votereg", "ideo5", "educ", "birthyr",
         "newsint"]]
         anes sub = anes[["ideo5","vote20cand"]]
         anes sub["vote20cand"] = anes_sub["vote20cand"].map(candidate_dict)
         /Users/quinnunderriner/anaconda3/lib/python3.7/site-packages/ipykernel
```

launcher.py:3: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame. Try using .loc[row indexer,col indexer] = value instead

See the caveats in the documentation: http://pandas.pydata.org/pandas-d ocs/stable/indexing.html#indexing-view-versus-copy

This is separate from the ipykernel package so we can avoid doing imp orts until

```
In [101]: #look at ideological spread
          anes_sub['ideo5'].value_counts()
Out[101]:
           3
                685
           4
                475
           2
                426
           1
                326
                302
           6
           5
                285
          -7
                  1
          Name: ideo5, dtype: int64
In [91]: #ready data for use in regression by making dummy variables of the ideol
          ogy figures and drop the non-responces
          spectrum dummies = pd.get dummies(anes sub, columns=['ideo5'])
          spectrum dummies = spectrum dummies.dropna(subset=['vote20cand'])
          spectrum dummies = pd.get dummies(spectrum dummies, columns=['vote20can
          d'])
          spectrum dummies = spectrum dummies[spectrum dummies["ideo5 -7"] != 1]
```

```
In [95]: import matplotlib.pyplot as plt
   import seaborn as sns
   import statsmodels.api as sm
   y = spectrum_dummies["vote20cand_Bernie Sanders"]
   X = spectrum_dummies[["ideo5_1","ideo5_2", "ideo5_3", "ideo5_4", "ideo5_
   5", "ideo5_6"]]
   X = sm.add_constant(X)

model = sm.OLS(y, X).fit()
   predictions = model.predict(X)
```

/Users/quinnunderriner/anaconda3/lib/python3.7/site-packages/numpy/core/fromnumeric.py:2389: FutureWarning: Method .ptp is deprecated and will be removed in a future version. Use numpy.ptp instead. return ptp(axis=axis, out=out, **kwargs)

Out[95]:

OLS Regression Results

Dep.	Variable	vote20	vote20cand_Bernie Sander			R-sq	uared:	0.012
	Model:	:		(OLS	Adj. R-squared:		0.008
	Method:	:	Le	ast Squ	ares	F-st	atistic:	2.887
	Date:	:	Mon, 0)9 Dec 2	019	Prob (F-statistic):		0.0134
	Time	1	14:12:13		2:13	Log-Likelihoo		-612.67
No. Obse	ervations	1		1	243		AIC:	1237.
Df R	esiduals:	:		1	237		BIC:	1268.
ı	Df Model:	:			5			
Covaria	nce Type:	:		nonrol	oust			
		-4-1		ъ и	[0.00	F 0.0751		
	coef	std err	t	P> t	[0.02	5 0.975]		
const	0.1782	0.014	13.122	0.000	0.15	2 0.205		
ideo5_1	0.0579	0.024	2.414	0.016	0.01	1 0.105		
ideo5_2	-0.0166	0.022	-0.745	0.456	-0.06	0.027		
ideo5_3	0.0088	0.022	0.404	0.686	-0.03	4 0.052		
ideo5_4	-0.0115	0.038	-0.304	0.761	-0.08	6 0.063		
ideo5_5	-0.0032	0.055	-0.058	0.954	-0.11	1 0.104		
ideo5_6	0.1428	0.040	3.598	0.000	0.06	5 0.221		
		000 004	.			0.000		
On	nnibus:	266.064	Durbir	n-Watso	n:	2.026		
Prob(Om	nibus):	0.000	Jarque-	Bera (JI	3):	466.953		
	Skew:	1.493		Prob(JI	3): 4.	00e-102		
Kı	ırtosis:	3.307	(Cond. N	o. 2	.59e+15		

Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The smallest eigenvalue is 2.31e-28. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

The two statistically significant predictor of support for Bernie Sanders is self-identification of being the "most liberal" on a 1-6 ideology scale, as well as the most conservative. Given that those listed are already likley democratic voters this is a really interesting finding. The coefficient on most liberal as a predictor is slighly higher for Bernie than Warren, but not by a significant amount

```
In [387]: y = spectrum_dummies["vote20cand_Elizabeth Warren"]
X = spectrum_dummies[["ideo5_1","ideo5_2", "ideo5_3", "ideo5_4", "ideo5_
5", "ideo5_6"]]
X = sm.add_constant(X)

model = sm.OLS(y, X).fit()
predictions = model.predict(X)

model.summary()
```

/Users/quinnunderriner/anaconda3/lib/python3.7/site-packages/numpy/cor e/fromnumeric.py:2389: FutureWarning: Method .ptp is deprecated and wil l be removed in a future version. Use numpy.ptp instead. return ptp(axis=axis, out=out, **kwargs)

Out[387]:

OLS Regression Results

Dep.	Variable:	vote20	vote20cand_Elizabeth Warren			R-s	quared:	0.007
	Model:				OLS	Adj. R-s	quared:	0.003
	Method:		l	_east Sc	luares	F-s	tatistic:	1.640
	Date:		Wed, 11 Dec 2019		2019	Prob (F-st	tatistic):	0.146
	Time:		17:24:44			Log-Like	elihood:	-168.86
No. Obse	ervations:				1243		AIC:	349.7
Df F	Residuals:				1237		BIC:	380.5
I	Df Model:				5			
Covaria	nce Type:			nonr	obust			
	coef	std err	t	P> t	[0.025	0.975]		
const	0.0597	0.010	6.282	0.000	0.041	0.078		
ideo5_1	0.0549	0.017	3.270	0.001	0.022	0.088		
ideo5_2	0.0211	0.016	1.350	0.177	-0.010	0.052		
ideo5_3	0.0208	0.015	1.363	0.173	-0.009	0.051		
ideo5_4	-0.0041	0.027	-0.156	0.876	-0.056	0.048		
ideo5_5	-0.0597	0.038	-1.558	0.120	-0.135	0.015		
ideo5_6	0.0267	0.028	0.962	0.336	-0.028	0.081		
On	nnibus:	743.814	Durbi	n-Watso	on:	1.938		
Prob(Om	nibus):	0.000	Jarque-	·Bera (J	B): 42	23.309		
	Skew:	2.957		Prob(J	В):	0.00		
Kı	ırtosis:	9.824		Cond. N	lo. 2.5	59e+15		

Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The smallest eigenvalue is 2.31e-28. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

```
In [390]: y = spectrum_dummies["vote20cand_Joe Biden"]
X = spectrum_dummies[["ideo5_1","ideo5_2", "ideo5_3", "ideo5_4", "ideo5_
5", "ideo5_6"]]
X = sm.add_constant(X)

model = sm.OLS(y, X).fit()
predictions = model.predict(X)
model.summary()
```

/Users/quinnunderriner/anaconda3/lib/python3.7/site-packages/numpy/cor e/fromnumeric.py:2389: FutureWarning: Method .ptp is deprecated and wil l be removed in a future version. Use numpy.ptp instead. return ptp(axis=axis, out=out, **kwargs)

Out[390]:

OLS Regression Results

Dep.	Variable:	vote20	vote20cand_Joe Biden			R-squared:		
	Model:			OLS	Adj. F	R-squared:	0.040	
	Method:		Least S	quares	ı	F-statistic:	11.37	
	Date:	We	ed, 11 De	c 2019	Prob (F	-statistic):	9.63e-11	
	Time:		17	7:24:57	Log-L	ikelihood:	-782.27	
No. Obse	ervations:			1243		AIC:	1577.	
Df R	lesiduals:			1237		BIC:	1607.	
ı	Of Model:			5				
Covaria	nce Type:	nonrobust						
	coef	std err	t	P> t	[0.025	0.975]		
	0.2879	0.016	18.497	0.000	0.257	0.318		
const								
ideo5_1	-0.1178	0.027	-4.283	0.000	-0.172	-0.064		
ideo5_2	-0.0094	0.026	-0.366	0.715	-0.060	0.041		
ideo5_3	0.1251	0.025	4.997	0.000	0.076	0.174		
ideo5_4	0.1454	0.043	3.348	0.001	0.060	0.231		
ideo5_5	0.0621	0.063	0.989	0.323	-0.061	0.185		
ideo5_6	0.0825	0.045	1.813	0.070	-0.007	0.172		
On	nnibus:	1766.083	Durb	in-Wats	on.	1.995		
Prob(Om	nibus):	0.000	Jarque	-Bera (JB): I	94.323		
	Skew:	0.736		Prob(JB): 6.3	36e-43		
Κι	ırtosis:	1.740		Cond.	No. 2.5	59e+15		

Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The smallest eigenvalue is 2.31e-28. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

For Joe Biden we see a wider spectrum. Buttigieg was not in the ANES data as he did not announce his candincy for president until April 14, 2019.

```
In [389]: y = spectrum_dummies["vote20cand_Amy Klobuchar"]
X = spectrum_dummies[["ideo5_1","ideo5_2", "ideo5_3", "ideo5_4", "ideo5_
5", "ideo5_6"]]
X = sm.add_constant(X)

model = sm.OLS(y, X).fit()
predictions = model.predict(X)
model.summary()
```

/Users/quinnunderriner/anaconda3/lib/python3.7/site-packages/numpy/cor e/fromnumeric.py:2389: FutureWarning: Method .ptp is deprecated and wil l be removed in a future version. Use numpy.ptp instead. return ptp(axis=axis, out=out, **kwargs)

Out[389]:

OLS Regression Results

Dep.	Variable	vote20c	ote20cand_Amy Klobucha			R-sq	uared:	0.001
	Model	:			OLS	Adj. R-squared:		-0.003
	Method	:	Le	ast Squ	ares	F-st	atistic:	0.2497
	Date	:	Wed,	11 Dec 2	2019	Prob (F-sta	itistic):	0.940
	Time	Time: 17:24:55			4:55	Log-Likel	ihood:	271.48
No. Obse	ervations	:		1	1243		AIC:	-531.0
Df R	Residuals	:		1	1237		BIC:	-500.2
I	Df Model	:			5			
Covaria	псе Туре	:		nonro	bust			
	coef	std err	t	P> t	[0.02	5 0.975]		
					_	_		
const	0.0334	0.007	5.005	0.000	0.020	0.046		
ideo5_1	-0.0021	0.012	-0.180	0.857	-0.025	5 0.021		
ideo5_2	0.0112	0.011	1.022	0.307	-0.010	0.033		
ideo5_3	0.0056	0.011	0.521	0.602	-0.01	5 0.027		
ideo5_4	0.0111	0.019	0.595	0.552	-0.02	5 0.048		
ideo5_5	-0.0084	0.027	-0.311	0.756	-0.06	0.044		
ideo5_6	0.0160	0.019	0.822	0.411	-0.022	2 0.054		
On	nnibus:	1174.615	Durh	oin-Wats	eon:	2.037		
Prob(Om	nibus):	0.000	Jarque	e-Bera (JB) : 2	26111.470		
	Skew:	4.727		Prob(JB):	0.00		
Κι	ırtosis:	23.367		Cond.	No.	2.59e+15		

Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The smallest eigenvalue is 2.31e-28. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

```
In [388]: y = spectrum_dummies["vote20cand_Kamala Harris"]
X = spectrum_dummies[["ideo5_1","ideo5_2", "ideo5_3", "ideo5_4", "ideo5_
5", "ideo5_6"]]
X = sm.add_constant(X)

model = sm.OLS(y, X).fit()
predictions = model.predict(X)
model.summary()
```

/Users/quinnunderriner/anaconda3/lib/python3.7/site-packages/numpy/cor e/fromnumeric.py:2389: FutureWarning: Method .ptp is deprecated and wil l be removed in a future version. Use numpy.ptp instead. return ptp(axis=axis, out=out, **kwargs)

Out[388]:

OLS Regression Results

Dep.	Variable:	vote20	cand_Ka	mala Ha	rris		R-squ	ared:	0.051
	Model:		OLS			Adj. R-squared:			0.047
	Method:		Lea	ıst Squa	res		F-sta	tistic:	13.37
	Date:		Wed, 1	1 Dec 20)19 	Prob ((F-stat	istic):	1.01e-12
	Time:			17:24	:54	Log-	-Likeli	hood:	-191.29
No. Obse	ervations:			12	243			AIC:	394.6
Df F	Residuals:			12	237			BIC:	425.3
!	Df Model:				5				
Covaria	nce Type:			nonrob	ust				
	coef	std err	t	P> t	[0.02	25 0	.975]		
const	0.0590	0.010	6.096	0.000	0.04	10 (0.078		
ideo5_1	0.1355	0.017	7.925	0.000	0.10)2 (0.169		
ideo5_2	0.0524	0.016	3.297	0.001	0.02	21 (0.084		
ideo5_3	-0.0226	0.016	-1.454	0.146	-0.05	53 (800.0		
ideo5_4	-0.0257	0.027	-0.950	0.342	-0.07	79 (0.027		
ideo5_5	-0.0340	0.039	-0.871	0.384	-0.11	11 (0.043		
ideo5_6	-0.0466	0.028	-1.650	0.099	-0.10	02 (0.009		
0	nnibus:	643.887	Dlai	\\/_+		1.0	961		
On	nnibus:	043.007	Durbi	n-Watso	on:	1.8	101		
Prob(Om	nibus):	0.000	Jarque-	Bera (J	B) : 2	2841.0	036		
	Skew:	2.587		Prob(J	B):	0	.00		
Kı	ırtosis:	8.299		Cond. N	lo. 2	2.59e+	-15		

Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The smallest eigenvalue is 2.31e-28. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

```
In [377]: y = spectrum_dummies["vote20cand_Beto O'Rourke"]
X = spectrum_dummies[["ideo5_1","ideo5_2", "ideo5_3", "ideo5_4", "ideo5_
5", "ideo5_6"]]
X = sm.add_constant(X)

model = sm.OLS(y, X).fit()
predictions = model.predict(X)
model.summary()
```

/Users/quinnunderriner/anaconda3/lib/python3.7/site-packages/numpy/cor e/fromnumeric.py:2389: FutureWarning: Method .ptp is deprecated and wil l be removed in a future version. Use numpy.ptp instead. return ptp(axis=axis, out=out, **kwargs)

Out[377]:

OLS Regression Results

Dep.	Variable:	vote20	vote20cand_Beto O'Rourke		ırke	R-squared:		uared:	0.037
	Model:			(DLS	Adj. R-squar		uared:	0.033
	Method:		Le	ast Squa	ares	F-statistic		atistic:	9.569
	Date:		Tue, 1	0 Dec 2	019	Pro	ob (F-sta	itistic):	5.69e-09
	Time:			20:55	5:25	L	.og-Likel	ihood:	-450.51
No. Obse	ervations:			1	243			AIC:	913.0
Df F	lesiduals:		1237					BIC:	943.8
I	Df Model:				5				
Covaria	nce Type:			nonrok	oust				
	coef	std err	t	P> t	[0.0	25	0.975]		
const	0.0862	0.012	7.235	0.000	0.0		0.110		
ideo5_1	0.1013	0.021	4.810	0.000	0.0	60	0.143		
ideo5_2	0.1338	0.020	6.832	0.000	0.0	95	0.172		
ideo5_3	0.0281	0.019	1.464	0.144	-0.0	10	0.066		
ideo5_4	-0.0418	0.033	-1.257	0.209	-0.10	07	0.023		
ideo5_5	-0.0612	0.048	-1.274	0.203	-0.1	56	0.033		
ideo5_6	-0.0739	0.035	-2.121	0.034	-0.1	42	-0.006		
On	nnibus: 3	396.037	Durhi	n-Watso	n.		1.938		
Prob(Om		0.000	Jarque-			87	75.984		
FIUD(UIII	Skew:	1.864	varque-	Prob(J	•		Se-191		
Κι	ırtosis:	4.738		Cond. N	•		9e+15		
						-	-		

Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The smallest eigenvalue is 2.31e-28. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

For someone more centrist like Klobuchar, none of the self-identified ideological points

```
In [359]: #look at total number of people across ANES data that said they would vo
    te for a given candidate
    anes_rank = pd.DataFrame(anes_sub.vote20cand.value_counts()).reset_index
    ()
```

In [360]: anes_rank

Out[360]:

	index	vote20cand
0	Joe Biden	391
1	Bernie Sanders	246
2	Beto O'Rourke	183
3	Kamala Harris	115
4	Elizabeth Warren	105
5	Cory Booker	54
6	Amy Klobuchar	49
7	Kirsten Gillibrand	38
8	Chris Murphy	26
9	Eric Holder	22
10	Deval Patrick	14

```
In [146]:
          def extact related words(df):
              grab words surrounding a candidates name to create corpus for each c
          andidate
              sanders list = []
              warren list = []
              biden list = []
              buttigieg_list = []
              beto_list = []
              harris list = []
              klobuchar_list = []
              for index, word in enumerate(df):
                  if word == "sander":
                       sanders_list.append(df[index-10:index+10])
                  if word == "warren":
                      warren_list.append(df[index-10:index+10])
                  if word == "biden":
                       biden list.append(df[index-10:index+10])
                  if word == "buttigieg":
                      buttigieg_list.append(df[index-10:index+10])
                  if word == "beto":
                      beto list.append(df[index-10:index+10])
                  if word == "harri":
                      harris list.append(df[index-10:index+10])
                  if word == "klobuchar":
                      klobuchar list.append(df[index-10:index+10])
              return sanders list, warren list, biden list, buttigieg list, beto l
          ist, harris list, klobuchar list
```

```
In [153]:
              #generate sentiment score for each candidate across all six months
              sanders_list_text = " ".join(map(str, sanders_list))
              warren_list_text = " ".join(map(str, warren_list))
              biden_list_text = " ".join(map(str, biden_list))
              buttigieg_list_text = " ".join(map(str, buttigieg_list))
              beto_list_text = " ".join(map(str, beto_list))
              harris_list_text = " ".join(map(str, harris_list))
              klobuchar_list_text = " ".join(map(str, klobuchar list))
              afinn = Afinn()
              biden_score = afinn.score(biden_list_text)
              sanders_score = afinn.score(sanders list text)
              buttigleg score = afinn.score(buttigleg list text)
              warren_score = afinn.score(warren_list_text)
              buttigieg_score = afinn.score(buttigieg_list_text)
              beto score = afinn.score(beto list text)
              harris score = afinn.score(harris list text)
              klobuchar_score = afinn.score(klobuchar_list_text)
              full_sentiment = pd.DataFrame(np.array([[biden_score],
                                             [sanders_score],
                                             [buttigieg_score],
                                             [warren_score],
                                                       [beto_score],[harris_score],
          [klobuchar score]]),
                                                       columns=["sentiment scores"
          ],
                                                       index=["Biden", "Sanders", "Bu
          ttigieg","Warren","beto","harris","klobuchar"])
```

```
In [154]: #generate rank for sentiment score accross the whole corpus
full_sentiment["sentiment_rank"] = full_sentiment["sentiment_scores"].ra
nk(ascending=False)
```

In [363]: full_sentiment.sort_values(by='sentiment_rank', ascending=True)

Out[363]:

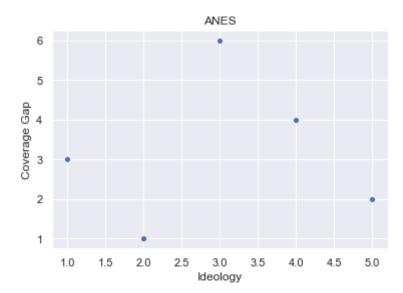
	index	sentiment_scores	sentiment_rank
3	Elizabeth Warren	2262	1
0	Joe Biden	1491	2
1	Bernie Sanders	1382	3
6	Amy Klobuchar	678	4
4	Beto O'Rourke	448	5
5	Kamala Harris	417	6
2	Pete Buttigieg	330	7

```
In [372]: #generate DF with candidate level difference between sentiment rank and
           ANES stated voting preference rank
          anes_rank = anes_rank[["index","vote20cand"]]
          full_comparison = anes_rank.merge(full_sentiment, on="index",how="inner"
          anes rank["Voting rank"] = anes rank["vote20cand"].rank(ascending=False)
          full_comparison = full_comparison[["index","Voting_rank","sentiment_ran
          k"]]
          full comparison["difference"] = full comparison["sentiment rank"] - full
          comparison["Voting rank"]
          #full comparison = full comparison[["index","vote20cand","sentiment scor
          es"]]
          full comparison = full comparison[["index", "sentiment rank", "Voting ran
          k", "difference" | |
          full comparison = full comparison.rename(columns={"index": "Candidate",
          "sentiment_rank": "Sentiment Rank", "Voting_rank": "Voting Rank", "differenc
          e": "Difference" })
          full comparison["Voting Rank"] = full comparison["Voting Rank"].astype(i
          nt)
```

```
In [458]: ideology2 = np.array([5,1,3,2,4])
    full_comparison['ideology'] = np.array(ideology2)
    ax = sns.scatterplot(x="ideology", y="sentiment_rank", data=full_comparison)

    plt.xlabel('Ideology')
    plt.ylabel('Coverage Gap')
    plt.title('ANES')
```

Out[458]: Text(0.5, 1.0, 'ANES')



```
In [238]: #ready text for use in topic model
    test_text = ""
    for i in sanders_list:
        for j in i:
            test_text += (j + " ")
    test_text = [test_text]
```

```
In [274]:
          def text topic model(df):
               make text right format for topic models
               new_list = []
               for i in df:
                   new_list.append(' '.join(i))
               return new list
In [259]: len(new list bernie)
Out[259]: 1465
In [247]: | from sklearn.feature_extraction import text
           stop words = text.ENGLISH STOP WORDS.union("ms","candid","mr","said")
In [281]:
               stop_words = set(stopwords.words('english'))
               #add specific stopwords plus names of candidates
               stop_words.update(["would", "said", "mr", "presid", "booker", "democrat",
           "new", "like", "elect", "year", "mrs", \
                                   "american", "also", "issu", "june", "moslty", "could",
           "whether", "day", "go", "use", "two", "one", "campaign", \
                                  "call", "even", "say", "get", "may", "make", "come", "cam
           paign","state","polit", \
                                   "time", "ms", "call", "elizabeth", "warren", "joe", "bid
           en","pete","buttigieg","bernie", \
                                  "sander", "berni", "rourk", "debat", "campaign", "candi
           d","first","second","percent","night","poll", \
                                  "vice", "vermont", "massachusett", "joseph", "kamala",
           "south", "bend", "jr", "senat", "harri", "race", "former", "th"])
```

```
In [273]: #topic modeling for bernie
          count vectorizer = CountVectorizer(stop words=stop words)
          count data bernie = count vectorizer.fit transform(new list bernie)
          def print topics(model, count vectorizer, n top words):
              words = count vectorizer.get feature names()
              for topic_idx, topic in enumerate(model.components_):
                  print("\nTopic #%d:" % topic_idx)
                  print(" ".join([words[i]
                                   for i in topic.argsort()[:-n_top_words - 1:-1
          ]]))
          number topics = 7
          number_words = 8
          # LDA model
          lda = LDA(n_components=number_topics, n_jobs=-1)
          lda.fit(count data bernie)
          print("Topics found via LDA:")
          print_topics(lda, count_vectorizer, number_words)
```

/Users/quinnunderriner/anaconda3/lib/python3.7/site-packages/sklearn/de composition/online lda.py:536: DeprecationWarning: The default value fo r 'learning method' will be changed from 'online' to 'batch' in the rel ease 0.20. This warning was introduced in 0.18. DeprecationWarning) Topics found via LDA: Topic #0: health care among lead nation medicar gener trump Topic #1: voter endors top right look creat team media lead show point larg els appear rival onstag Topic #3: progress includ sever class much expect york support Topic #4: plan presidenti money repres million rais vote tax Topic #5: side stand left polici white mostli seek nomin Topic #6: donor small individu far mayor moder andrew estim

```
In [276]: #generate list for biden
          new list biden = text topic model(biden list)
In [454]: count vectorizer = CountVectorizer(stop words=stop words)
          count data biden = count vectorizer.fit transform(new list biden)
          def print topics(model, count vectorizer, n top words):
              words = count_vectorizer.get_feature names()
              for topic_idx, topic in enumerate(model.components_):
                  print("\nTopic #%d:" % topic idx)
                  print(" ".join([words[i]
                                   for i in topic.argsort()[:-n_top_words - 1:-1
          11))
          number_topics = 3
          number words = 8
          # LDA model
          lda = LDA(n components=number topics, n jobs=-1)
          lda.fit(count data biden)
          print("Topics found via LDA:")
          print_topics(lda, count_vectorizer, number_words)
          /Users/quinnunderriner/anaconda3/lib/python3.7/site-packages/sklearn/de
          composition/online lda.py:536: DeprecationWarning: The default value fo
          r 'learning method' will be changed from 'online' to 'batch' in the rel
          ease 0.20. This warning was introduced in 0.18.
            DeprecationWarning)
          Topics found via LDA:
          Topic #0:
          trump voter support clinton front lead primari iowa
          trump investig son hunter ukrain ukrainian zelenski corrupt
          Topic #2:
          trump rais fund far donor back attack polici
```